Post Your Question (Due November 14th)

Subscribe

Start a New Thread

Filter by: All Threads

Sort by: Least Recent Activity

Companionship AI

Marcelo Figueroa Parada posted Nov 14, 2025 1:00 AM Subscribe

Introduction

Some of the articles I read discuss the reasons and measurable benefits certain people can get for using companionship AI. They have potential to aid people going through break ups, practice therapeutic exercises outside of therapy sessions, cope with loneliness, etc. One article specifically sets a 2-sided framework to determine if companionship AI is problematic or not: Substitute vs Supplement. If a companion AI supplements your social life or someone else's social role (like a therapist or teacher), it can be good or at least generally unproblematic. But if an AI substitutes someone's social role then many dangers arise: can hurt how we relate with others, may entrench the problems that the AI was supposed to help with (like worsen your isolation), incentivize gross business practices (if we depend on the AI for emotional needs, then companies can monetize that), can hurt those around us if we replace their social roles, and pose greater societal dangers if certain demographic segments use companion AI's over participating in social roles.

Question

I'm not entirely convinced that pushing away from companion AIs that pose the risk of substituting social roles is a good approach. Other highly regulated products and services like drugs, sex work, and gambling have a complicated history with regulation. Banning them can increase their illicit use and dangerous practices like laced drugs, human sex trafficking, unregulated gambling rings, etc. Bans can also complicate victims access to help and stunt research into more ethical uses. Do you think that regulating against the substitution risks of companion AI aligns with ethical practice, even with the risks that come with illicit use? Or should we allow substitution uses to be developed and center our ethical concerns elsewhere when regulating companionship AI?

less

0 Unread 0 Replies 4 Views

From Fear to Fairness: What Do We Genuinely Need from an AI We Can Trust?

Piter Garcia posted Nov 14, 2025 1:23 AM • **502** Words Subscribe

Hi everyone,

Something I’ve come to understand through my research, and my own life, is that most conversations about AI begin from a place of **fear**. And that’s okay. Fear is a natural response to powerful, unfamiliar technology. I’ve certainly felt it. But I've also learned that if we only ever

react to our fears, we risk building solutions that are fast and direct, but that **never fix the root causes** of what we're afraid of, ultimately limiting our understanding. I believe fear is a phase we must **transition *through*, not a place to stay**. Sometimes, the thing we fear the most, if we can approach it with a "**safety net**" of ethical principles, can become the very thing that helps us the most.
My research portfolio is built on this idea. It focuses on creating **equitable AI** that marginalized communities, especially CLD (Culturally and Linguistically Diverse) communities, can actually **trust**. The core of my EQUITAS framework, informed by both my research and personal experience, is that genuine **community engagement** is the only way to build that trust. We can't just build AI *for* people; we have to build it ***with*** them, especially now, during what I see as the **birth and childhood of AI**.

This has led me to a fundamental question that goes beyond just mitigating bias. Since fairness and equity will inevitably look different from one community to the next, I genuinely want to ask you all to reflect on your own perspectives.

1. **What does "fairness" in an AI actually look like to you personally?** Not in the abstract, but for you. What would an AI need to do or be for you to consider it fair?
    - *For example, for a student applying to college, does fairness mean the AI ignores their zip code and socioeconomic background completely (equal treatment)? Or does it mean the AI takes that context into account to level the playing field (equal opportunity)?*
2. Now, consider that personal definition of fairness. **Can you imagine a way that this version of fairness, if applied universally, might inadvertently cause harm or feel *un*fair to another person or group within your own community?**
    - *For instance, if an AI is designed to be "fair" by giving every patient exactly 15 minutes of a doctor's time, would that feel unfair to an elderly patient who speaks English as a second language and needs more time to communicate, compared to a young, tech-savvy patient who can get their point across in five minutes?*
3. If so, how do we resolve that tension? **How can we design an AI that is fair to *you* without perpetuating harm to others?**
    - *Continuing the example, could the AI be fair by allocating time based not on equality (15 minutes for all), but on equity (allocating more time to those who need it)? What new problems might that create?*

I'm trying to shift the conversation from "what are you afraid of?" to "what do you truly need?" and then figure out how we can build that together.

less

0	0	4
Unread	Replies	Views

AI Rights Question

Dan Dumouchel posted Nov 14, 2025 1:43 PM Subscribe

Portfolio Summary

My portfolio establishes a clear distinction between the philosophical threshold for AI rights and the practical trigger for immediate policy. While I agree with philosophers like DeGrazia, Andreotta, and Searle that true moral status must be grounded in, I find their conclusion of inaction unpersuasive because of the epistemic problem. We can never know for certain if a non-biological AI is conscious.

Therefore, my core belief centers on a Moral Precautionary Principle. I argue that we must use functional complexity, like the ability to set complex goals (DeGrazia's counterpoint) or trigger strong human social bonding (Mollen et al., Darling), as an immediate, extrinsic trigger for governance. This ensures that we act responsibly now to mitigate both the societal risks (Bostrom, Sharma) and the moral risk of abusing a potentially conscious entity.

My Questions

The Epistemic Problem (Andreotta, Searle) demands that we cannot know if AI has intrinsic moral status (DeGrazia's sentience criterion). Therefore, should our immediate ethical priority be to mandate agile governance (Wallach) and extrinsic protections (Bryson, Sharma) based on functional complexity (Mollen et al.)?

What is the long-term ethical cost of permanently institutionalizing a system where an entity's treatment is governed by its function (a tool standard) rather than its potential intrinsic nature (a

moral patient standard)?

less

0 Unread | 0 Replies | 1 Views

On AI Rights and Human Exceptionalism

Zane Draper posted Nov 14, 2025 1:43 PM Subscribe

In the realm of animal ethics, humans are often argued to hold a higher moral status based on a higher level of perceived sentience. This would seem to indicate sentience comes as a scale, and that the more sentient an entity is, the more moral worth it has. And yet, the most common counter argument is that we do not consider those of less capability within our society, like babies or those of with intellectual disabilities, to be of less moral worth. This would seem to indicate that scale of sentience is not the dividing line; human exceptionalism is. We perceive ourselves as more deserving of moral rights simply because we are human.

Then when considering AI's right to moral status or ethical consideration, the main question to overcome is this: given an AI of sufficiently comparable intellectual and emotional complexity to a human, how do we overcome the perceive bias we have towards human exceptionalism and biological origins? How can a artificial entity truly be argued to be equal to that of a biological one?

less

0 Unread | 0 Replies | 2 Views

AI for Autonomous Weapon Systems

Ryan Kenward posted Nov 14, 2025 3:53 PM Subscribe

My research focuses on Autonomous Weapon Systems (AWS) and Lethal Autonomous Weapon Systems (LAWS). A recurring challenge in the literature is the lack of consistent definitions for these systems and the concept of *autonomy* itself. Public perception often shaped by science fiction (e.g., *Terminator*) diverges significantly from technical reality.

In particular:

- Philosophical discussions of autonomy (e.g., moral agency, free will) dominate ethical debates.
- Engineering definitions, by contrast, describe autonomy pragmatically: a self-contained system operating within predefined rules and without real-time human intervention.

This disconnect fuels hype and fear, overshadowing legitimate national security needs—such as automated defense against high-volume threats (e.g., Israel's Iron Dome responding to missile barrages).

I'd like to hear perspectives on:

1. How to establish clearer, technically grounded definitions of AWS/LAWS and autonomy.
2. Strategies to redirect public and policy discourse away from dystopian narratives toward pragmatic risk-benefit analysis.
3. Ways to emphasize that autonomy != self-awareness and that robust rule-based constraints can bound system behavior.

Ideas for framing these systems as constrained, predictable tools—rather than emergent intelligences—would be especially welcome.

less

0 Unread | 0 Replies | 2 Views

Caring and Caregiving in the age of AI

Rachel Thompson posted Nov 14, 2025 9:30 PM Subscribe

I initially started thinking about two, perhaps somewhat disparate things - first, AI systems within the caregiving field - in particular, systems that help assist the care for people with disabilities and older adults (for example, companion robots, care bots, and AI disability support agents and assistive technology). At the same time, I started thinking about something I often see in my day to day work as a solutions lead at a policy institute in which we see* the creation and deployment of AI systems specifically designed to ration or take away care or minimize the cost and burden of care (see also Eubanks). At first these two lines of thinking seemed different, but eventually I started to ask myself a more fundamental set of questions - what would it mean for AI systems to care or caregive? Can they? And what exactly does it mean to give care in the first place? What would a future look like in which AI participates in caregiving and works with us to fulfill and amplify (rather than to deny and erode) the human duty of care to one another? How might such systems also help us distribute the labor of care more equitably?

To clarify this idea, I draw on an analogy with service animals and therapy animals. These animals are not moral agents: they do not possess moral autonomy, reflective empathy, or a conceptual understanding of duty. Yet they participate meaningfully and authentically in caregiving networks. They do so not because of internal moral reasoning, but because care ethics is fundamentally relational and functional (drawing on relational framework from Birhane). A service dog supports care through attentiveness, responsiveness, competence, and relational integration - qualities that embed it in a caregiving practice and give it a role within institutions of care. Importantly, the dog is not a mere tool; its participation does not absolve humans of responsibility, but becomes part of a shared practice of caring.

This analogy makes me wonder whether AI systems might someday participate in care networks in a similar way, not as moral agents, but as relational partners shaped by human values, norms, and duties of care. If so, what capacities or design principles would such AI systems need in order to count as authentic participants in caregiving, rather than as tools that merely simulate care or enable institutions to offload responsibility?

Obviously, this is kind of a lot of different thoughts, so I could use some help with a few questions:

- **If care ethics centers relational qualities rather than internal emotional states (Tronto, Held, Birhane), what criteria should determine whether an AI system can legitimately participate in a caregiving relationship?**
- **A major objection in the literature (for example, Vallor) is that machines without feelings can only simulate care (and maybe only in ethically problematic ways). How should I approach this objection, and is the service animal analogy a plausible way to respond to it? And actually does the service animal analogy really belong in my paper?**
- **I think we often think about caregiving in a sort of 'warm and fuzzy' way, but as a 'sandwich generation' member firmly squashed between caring for my child and my aging parents, as well as responsible for mentorship and growth of dozens of junior colleagues at work, I do NOT feel the warm and fuzzy these days! Caregiving feels very onerous, often unwanted, and largely uncompensated. I think often the objections to technological caregiving may come from a place of the 'warm and fuzzy', generally idealistic and not rooted in the corporeal reality of what it truly means to care for others. Is it possible that some of the objections to technologically mediated caregiving come from a place of idealism and do not fully grapple with the reality of caregiving or the unequal burden it places on some groups of people (for example, specifically women and people of color within a**

certain age bracket)? How might AI caregiving not only even out these types of burdens, but also itself provide much needed care to caregivers? For example, if I set my aging parent up with a smartwatch that reminds her to take meds and monitors her for falls, that is actually care *for me* because I don't need to go over there every five seconds.**

- **Finally, for the purposes of narrowing this into a focused research paper, what specific example or technology should I anchor my analysis in - caregiving robots, AI companionship tools, assistive technology, or something else?**

* I do not and will not work on such systems, and nor does the company I work for. But we often see reference to them in the form of client communications and RFPs that ask contractors to construct such tools. Most recently, I saw an RFP from the federal government asking for contractors who would construct an AI system to deny what it deemed as 'unnecessary' health care for Medicaid participants. As anyone in healthcare policy would tell you, there most certainly is such a thing as non-evidence based health care that perhaps as a society we should consider no longer providing; but that is very different than the current administration deciding, using black-box AI tools, that particular people do not get to receive particular types of medical care.

** As you can see the other underlying factor behind why I chose this topic is that I am TIRED of caregiving and TESTY :)

less

0	0	2
Unread	Replies	Views

Question – Aaliyan Qadri

Aaliyan Qadri posted Nov 11, 2025 12:34 PM Subscribe

Brief of My Research

My research examines how major technology corporations defend their participation in military artificial intelligence projects. These companies often claim that their technologies enhance safety, precision, and national security. However, I argue that these justifications often produce ethical distortion, where ethics is used as an external defense rather than as an internal moral commitment. In other words, corporations may appear ethical while using moral reasoning as a strategic tool to protect profit and political interests instead of upholding genuine accountability.

Discussion Question

If a corporation genuinely believes that its military artificial intelligence work promotes safety or national security, but that belief also serves its financial and political goals, can its ethical reasoning still be considered authentic? At what point does moral justification turn into ethical distortion, and how can we tell when a company's ethical claims are sincere rather than self serving?

less

1	1	4
Unread	Replies	Views

Last post **1 hour ago** by Marcelo Figueroa Parada

Start a New Thread